

CS 839 HW1

Ivy Zhu

September 2025

1 GPT-2 parameter counts

1. Write symbolic formulas for embeddings, attention, MLP, and LayerNorms. Give a total in terms of V , E , H , L , P , where: V = vocabulary size, E = model/embedding dimension, H = number of attention heads, L = number of transformer layers, and P = maximum positional indices (context length). [8 pts]

- Embeddings = $wte + wpe = V * E + P * E$
wte: lookup the embedding of the input tokens in a $V * E$ matrix.
wtp: lookup the embedding of the positions of the tokens in a $P * E$ matrix.
- Attention = $c_attn + c_proj + attn_dropout + resid_dropout = 3E^2 + 3E + E^2 + E + 0 + 0 = 4E^2 + 4E$
 - $c_attn = E \times 3E + 1 \times 3E = 3E^2 + 3E$
This layer projects the normalized embeddings from ln_1 to Q, K, V.
It's $y = xW^T + b$, where x is the input embedding from the last layer ($1 \times E$), and W^T ($E \times 3E$) and b ($1 \times 3E$) are trainable.
 - $c_proj = E \times E + E = E^2 + E$
A Conv1D layer that combines the output of attention heads.
Multi-head attention splits up the embedding dimension, not duplicate it, so the input dimension here is still E . W^T is a $E \times E$ matrix and b is a $1 \times E$ vector.
 - $attn_dropout = 0$
A dropout layer, nothing learnable
 - $resid_dropout = 0$
A dropout layer, nothing learnable
- MLP = $c_fc + c_proj + act + dropout = 4E^2 + 4E + 4E^2 + E + 0 + 0 = 8E^2 + 5E$
"Feed-forward layer" or "multi-layer perceptron", applied to each token separately.

- $c_fc = E \times 4E + 4E = 4E^2 + 4E$
 "Up-projecting": turn E into $4E$. So it's a matrix of $E \times 4E$.
 Plus a bias with size $4E$.
 - $c_proj = 4E \times E + E = 4E^2 + E$
 "Down-projecting": turn it back into the dimension of embeddings. So it's a matrix of $4E \times E$. Plus a bias with size E .
 - $act = 0$
 Apply NewGELUActivation, nothing to learn.
 - $dropout = 0$
 Dropout layer, nothing to learn.
 - LayerNorms (within each GPT2Block) = $ln_1 + ln_2 = 2 * E + 2 * E = 4 * E$
 There's also a final LayerNorm layer (ln_f), which is also $2 * E$.
 Normalize the input across the embedding dimension. It's calculated for each embedding by $y = \frac{x - E[x]}{\sqrt{Var[x] + \epsilon}} * \gamma + \beta$, and the only trainable parameters are γ and β , which each has the same dimension as the embedding (E).
 - Total = $V * E + P * E + L * (12E^2 + 13E) + 2E$
2. *Plug in the hyperparameters for GPT-2. Report totals and a breakdown by component as concrete parameter counts (numbers). Verify with code and explain any mismatch. [6 pts]*

Hyperparameters:

- E (embedding dimension) = 768
- V (vocabulary size) = 50257
- H (number of attention heads) = 12
- L (number of transformer layers) = 12
- P (maximum positional indexes) = 1024

Per-component parameter counts:

- Embeddings = $(V + P)E = (50257 + 1024) * 768 = 39,383,808$
- LayerNorms = $4E = 4 * 50257 = 3,072$
- Attention = $4E^2 + 4E = 2,362,368$
- MLP = $8E^2 + 5E = 4,722,432$

Total per layer = LayerNorms + MLP + Attention = 7,087,872

Total = Embeddings + $L * (\text{Total per layer}) + ln_f = 124,439,808$

Counting programatically (in `gpt2.py count_params()`) gives the same result.

3. *Estimate totals for GPT-2 Medium, GPT-2 Large, and GPT-2 XL using documented (E , H , L). Briefly explain which terms dominate scaling.*

- GPT-2 Medium: $L = 24$, $E = 1024$, $H = 16$.
Total = Embeddings + $L * (\text{Total per layer}) + \ln f = 354,823,168$
- GPT-2 Large: $L = 36$, $E = 1280$, $H = 20$.
Total = Embeddings + $L * (\text{Total per layer}) + \ln f = 774,030,080$
- GPT-2 XL: $L = 48$, $E = 1600$, $H = 25$.
Total = Embeddings + $L * (\text{Total per layer}) + \ln f = 1,557,611,200$

Scaling explanation:

- H (number of attention heads) doesn't change the number of parameters because it's only used to split up the embeddings.
- E (embedding dimension) scales up fast because each layer scales up at $O(E^2)$.
- L (number of transformer layers) also scales up by linearly scaling up $O(L)$ after the $O(E^2)$.

2 Modern model study

1. *Report the exact configuration you use (cite the source). [4 pts]*

gpt-oss-20b

huggingface

Model Card

Blog Posts that helped understanding the design and architecture:

- <https://cameronrwolfe.substack.com/p/gpt-oss>
- From GPT-2 to gpt-oss: Analyzing the Architectural Advances

Configuration

2. *Derive component-wise formulas (as in Problem 1) consistent with the chosen architecture, and compute both per-layer and total parameter counts using your conventions. [20 pts]*

Hyperparameters:

- L (number of transformer layers): 24
- E (embedding dimension) = 2,880
- V (vocabulary size) = 201,088
- $head_dim$ (Key/Query/Value head dimension) = 64
- num_q_heads (number of query heads) = 64

- num_kv_groups (number of key-value groups) = 8
- num_MoE_blocks (number of MoE blocks) = 32
- $experts_per_token = 4$

Parameter counts per component:

- Embedding = $V * E$
- Transformer layer = $E + E * head_dim * num_q_heads + head_dim * num_q_heads + 2 * (E * head_dim * num_kv_groups + head_dim * num_kv_groups) + num_q_heads + E + E * num_MoE_blocks + num_MoE_blocks + num_MoE_blocks * (3E^2 + 3E)$
 - Input RMSNorm = E
 - Grouped Query Attention = $E * head_dim * num_q_heads + head_dim * num_q_heads + 2 * (E * head_dim * num_kv_groups + head_dim * num_kv_groups) + num_q_heads$
 - * Q projection = $E * head_dim * num_q_heads + head_dim * num_q_heads$
Project dimension of E to dimension of Query Dimension * Number of query heads.
So there's a matrix of $E * head_dim * num_q_heads$ + a vector of $head_dim * num_q_heads$ (bias)
 - * K projection = $E * head_dim * num_kv_groups + head_dim * num_kv_groups$
Project dimension of E to dimension of Query Dimension * Number of key-value groups.
So there's a matrix of $E * head_dim * num_kv_groups$ + a vector of $head_dim * num_kv_groups$ (bias)
 - * V projection = $E * head_dim * num_kv_groups + head_dim * num_kv_groups$
Same as K
 - * Output linear projection = $head_dim * num_q_heads * E + E$
Project the attention heads back onto the dimension of embeddings.
 - * Learned bias for each attention head = num_q_heads
According to the model card, "Each attention head has a learned bias in the denominator of the softmax, similar to off-by-one attention and attention sinks, which enables the attention mechanism to pay no attention to any tokens".
 - Post attention RMSNorm = E
 - MLP/MoE = $E * num_MoE_blocks + num_MoE_blocks + num_MoE_blocks * (3E^2 + 3E)$
 - * router = $E * num_MoE_blocks + num_MoE_blocks$
The router selects a small set of feed forward modules for each token.
gpt-oss uses a standard linear router projection.

* num_MoE_blocks * Each Feed Forward
 For each Feed Forward: $3E^2 + 3E$ parameters
 · Linear Layer: $E * E + E = E^2 + E$ parameters
 There are three linear layers.
 Not all are always active.

- Final RMSNorm = E
 RoPE is applied on every transformer layer, not only on the input.
 However, unlike the positional embedding in GPT2, RoPE is not learned, so there's no parameters in RoPE.

Calculation:

- Token Embeddings: 579,133,440
- Attention: 26,550,144
 - Q_proj: 11,800,576
 - K_proj: 1,475,072
 - V_proj: 1,475,072
 - Out_proj: 11,799,360
 - Bias per head: 64
- MLP: 796,631,072
 - Router: 92,192
 - Each expert/feed forward: 24,891,840
- each RMSNorm: 2,880

Per layer parameter count: 823,186,976

Total = Embedding + L * Per layer count + final RMSNorm = 20,335,622,208

Total = $V * E + L * (E + E * head_dim * num_q_heads + head_dim * num_q_heads + 2 * (E * head_dim * num_kv_groups + head_dim * num_kv_groups) + num_q_heads + E + E * num_MoE_blocks + num_MoE_blocks + num_MoE_blocks * (3E^2 + 3E) + E$

Calculation code is submitted in gpt-oss.py

The total calculation doesn't include unembed because it's not in the configuration. However, unembed is included in the model card, which causes the difference in the total.

3. Report the final totals as concrete parameter counts (numbers). If MoE, also report active parameters per token. [12 pts]

Total (without unembed) = 20,335,622,208

Total (with unembed) = 20,914,757,184

Active Parameters per token: only 4 expert per token.

- active_MLP = router + expert_per_token * each_expert = 99,659,552

- $\text{active_per_layer} = \text{RMSNorm} + \text{Attention} + \text{RMSNorm} + \text{active_MLP}$
 $= 126,215,456$
 - $\text{active_total} = \text{embed} + L * \text{active_per_layer} + \text{RMSNorm} = 3,608,307,264$
4. *Provide code that demonstrates a programmatic check. [12 pts]*
 Programmatic check is submitted in gpt-oss.py
5. *Identify the key components that differ from GPT-2. For each:*
- *Summarize the design motivation in your own words. [12 pts]*
 - *Provide evidence of effectiveness from the paper/tech report (ablation, benchmark, or efficiency claim). [12 pts]*

Key component differences:

- (a) Positional embedding: GPT-2 uses a learnable positional embedding layer, while GPT-oss uses RoPE, which doesn't have any learnable parameters, and supports bigger context length. It provides a bigger context length, reduces the number of learnable parameters, and also supports relative positions.
 Evidence: longer context length. GPT-2 has 1024 maximum positional indexes and needs to learn the parameters for positional embedding, while GPT-oss has 131,072 maximum position embeddings without the need of training it.
- (b) Normalization layers: GPT-2 uses LayerNorm, while GPT-oss uses RMSNorm (root mean square normalization), which reduces the cost in computing mean and variance. RMSNorm is becoming popular recently because of its lower computational cost.
 Evidence: According to Root Mean Square Layer Normalization, "RMSNorm yields comparable performance against LayerNorm but shows superiority in terms of running speed with a speed-up of 7%-64%."
- (c) Attention: GPT-2 uses multihead attention, while GPT-oss uses grouped query attention, which reduced the number of key-value pairs, and also reduced the number of parameters required to calculate keys and values.
 Evidence: It reduced the number of parameters of projecting V and K from 11,800,576 (if use multi-head) to 1,475,072.
 Also, unlike GPT-2, which keeps the total dimension of all query heads the same as the dimension of embeddings, GPT-oss expanded the total query dimension to $64 * 64 = 4,096$, which is bigger than the dimension of embedding (2,880). The motivation might be to provide bigger exposure to the attention.
 Evidence: <https://arxiv.org/html/2508.16700v2>, compared to dense baselines, GPT-oss-20B has higher throughput and lower energy consumption.

- (d) Feed forward: GPT-2 contains a simple feed forward module, while GPT-oss-20B has a MoE module that contains 32 individual experts, with a router layer that decides which experts to be used for each token. The MoE architecture increases the performance by including a large number of experts, while keeping the computation low by only activate a small subset of them for each token.

Evidence: OpenAI claims that "Each model is a Transformer which leverages mixture-of-experts (MoE) to reduce the number of active parameters needed to process input".

6. *Give a short overall summary (bullets or a short paragraph) of the main design trends you observe. [8 pts]*

- Overall, the architecture design has remained very similar, with the same types of basic components (embedding, normalization layers, attention, and MLP), but each with improvements.
- The number of layers has decreased, while the complexity of each layer has increased.
- The size and portion of MLP has significantly increased, but there's also effort to keep computational cost low while retain the advantage of dense MLP.
- There's effort on changing individual modules to reduce computational cost or to make it more parallel-friendly, and removing modules that has been proving to have worse effect than expected.
- Context length has increased significantly.